

# AIML CAPSTONE PROJECT

**PROJECT NAME : CAPSTONE PROJECT - NLP**

**GROUP NAME : GROUP 12-NLP2-CAPSTONE-DEC22A**

**GROUP MEMBERS : DEVESH PAREEK**

**NAVITHA G**

**PRASHANTAGARWAL**

**SUKRITI MODGIL**

**VIJI**

**VISHWANATH SHEBINAKATTI**

TABLE OF CONTENTS

HEADINGS	PAGE NUMBER
INTRODUCTION	4
DATASET DESCRIPTION	5
FINDING DUPLICATE VALUES	9
DROPPING DUPLICATE VALUES	9
SUMMARY OF THE APPROACH TO EDA AND PRE-PROCESSING	10
DESIGN	15
TRAINING	18
TESTING	

## INTRODUCTION :

Machine translation (MT) is the process of translating text in one language to another language with the use of software by incorporating both computational and linguistic knowledge.

NMT is a recently formulated method for automatic translation with the help of deep neural networks. NMT uses a single large neural network for training. This structure comprised of encoder and decoder networks where the encoder consumes the input sentences to produce a vector representation, and the decoder takes this vector and outputs the target language words. Generally, both encoder and decoder networks are designed using the recurrent neural networks (RNN) or long short-term memory (LSTM) or gated recurrent unit (GRU) or bidirectional RNN, which are the alternatives to RNN. Even though RNN, especially LSTM, is theoretically proven for handling long-term dependencies in the sentences

Here, in this case study we use a simple RNN model combined with lstm nodes for our translation task.

# 1. Dataset Description

The data contains three pairs,

europarl-v7\_en\_de.txt file - 1920209 text

europarl-v7\_de\_en.txt - 1920209 text

commoncrawl\_en\_de.txt - 2399123 text

commoncrawl\_de\_en.txt – 2399123 text

news-commentary-v9\_en\_de.txt – 201854 text

news-commentary-v9\_de\_en.txt – 201854 text

```
In [39]: file_cc_de_en=open('commoncrawl_de_en.txt',encoding='Utf-8')
cc_de_en=file_cc_de_en.readlines()
file_cc_en_de=open('commoncrawl_en_de.txt',encoding='Utf-8')
cc_en_de=file_cc_en_de.readlines()
cc_de_en_df=pd.DataFrame(cc_de_en,columns=['German'])
cc_en_de_df=pd.DataFrame(cc_en_de,columns=['English'])
```

```
In [40]: file_ep_de_en=open('europarl-v7_de_en.txt',encoding='Utf-8')
ep_de_en=file_ep_de_en.readlines()
file_ep_en_de=open('europarl-v7_en_de.txt',encoding='Utf-8')
ep_en_de=file_ep_en_de.readlines()
ep_de_en_df=pd.DataFrame(ep_de_en,columns=['German'])
ep_en_de_df=pd.DataFrame(ep_en_de,columns=['English'])
```

```
In [41]: file_news_de_en=open('news-commentary-v9_de_en.txt',encoding='Utf-8')
news_de_en=file_news_de_en.readlines()
file_news_en_de=open('news-commentary-v9_en_de.txt',encoding='Utf-8')
news_en_de=file_news_en_de.readlines()
news_de_en_df=pd.DataFrame(news_de_en,columns=['German'])
news_en_de_df=pd.DataFrame(news_en_de,columns=['English'])
```

```
In [42]: print(cc_de_en_df.count())  
print(cc_en_de_df.count())  
print(ep_de_en_df.count())  
print(ep_en_de_df.count())  
print(news_de_en_df.count())  
print(news_en_de_df.count())
```

```
German      2399123  
dtype: int64  
English     2399123  
dtype: int64  
German      1920209  
dtype: int64  
English     1920209  
dtype: int64  
German       201854  
dtype: int64  
English      201995  
dtype: int64
```

All the 6 files we have both german and english content alternatively. These are our datasets by using this information we going to split up english and german language.

## Merging all the three datasets :

```
In [43]: cc_de_en_df.join(cc_en_de_df)
pd.merge(cc_de_en_df, cc_en_de_df, left_index=True, right_index=True)
cc_df = pd.concat([cc_de_en_df, cc_en_de_df], axis=1)
```

```
In [44]: cc_df.head()
```

```
Out[44]:
```

	German	English
0	iron cement ist eine gebrauchsfertige Paste, ...	iron cement is a ready for use paste which is ...
1	Nach der Aushärtung schützt iron cement die Ko...	iron cement protects the ingot against the hot...
2	feuerfester Reparaturkitt für Feuerungsanlagen...	a fire resistant repair cement for fire places, ...
3	Der Bau und die Reparatur der Autostraßen...\n	Construction and repair of highways and...\n
4	die Mitteilungen sollen den geschäftlichen kom...	An announcement must be commercial character.\n

```
In [45]: cc_df.count()
```

```
Out[45]: German      2399123
English      2399123
dtype: int64
```

```
In [46]: ep_de_en_df.join(ep_en_de_df)
pd.merge(ep_de_en_df, ep_en_de_df, left_index=True, right_index=True)
ep_df = pd.concat([ep_de_en_df, ep_en_de_df], axis=1)
ep_df.head()
```

```
Out[46]:
```

	German	English
0	Wiederaufnahme der Sitzungsperiode\n	Resumption of the session\n
1	Ich erkläre die am Freitag, dem 17. Dezember u...	I declare resumed the session of the European ...
2	Wie Sie feststellen konnten, ist der gefürchte...	Although, as you will have seen, the dreaded '...
3	Im Parlament besteht der Wunsch nach einer Aus...	You have requested a debate on this subject in...
4	Heute möchte ich Sie bitten - das ist auch der...	In the meantime, I should like to observe a mi...

```
In [47]: ep_df.count()
```

```
Out[47]: German      1920209
English      1920209
dtype: int64
```

# Concatenated all the dataframes into a single dataframe

```
In [48]: news_de_en_df.join(news_en_de_df)
pd.merge(news_de_en_df, news_en_de_df, left_index=True, right_index=True)
news_df = pd.concat([news_de_en_df, news_en_de_df], axis=1)
news_df.head()
```

```
Out[48]:
```

	German	English
0	Steigt Gold auf 10.000 Dollar?\n	\$10,000 Gold?\n
1	SAN FRANCISCO – Es war noch nie leicht, ein ra...	SAN FRANCISCO – It has never been easy to have...
2	In letzter Zeit allerdings ist dies schwierige...	Lately, with gold prices up more than 300% ove...
3	Erst letzten Dezember verfassten meine Kollege...	Just last December, fellow economists Martin F...
4	Und es kam, wie es kommen musste.\n	Wouldn't you know it?\n

```
In [49]: news_df.count()
```

```
Out[49]: German      201854
English    201995
dtype: int64
```

```
In [14]: final_data_df = pd.concat([news_df, ep_df, cc_df])
```

```
In [15]: del[news_df, ep_df, cc_df]
```

```
In [16]: final_data_df.head()
```

```
Out[16]:
```

	German	English
0	Steigt Gold auf 10.000 Dollar?\n	\$10,000 Gold?\n
1	SAN FRANCISCO – Es war noch nie leicht, ein ra...	SAN FRANCISCO – It has never been easy to have...
2	In letzter Zeit allerdings ist dies schwierige...	Lately, with gold prices up more than 300% ove...
3	Erst letzten Dezember verfassten meine Kollege...	Just last December, fellow economists Martin F...
4	Und es kam, wie es kommen musste.\n	Wouldn't you know it?\n

```
In [17]: print(final_data_df.count())
```

```
German      4521186
English    4521327
dtype: int64
```



## FINDING DUPLICATE VALUES

```
In [19]: final_data_df[final_data_df.duplicated()]
```

```
Out[19]:
```

	German	English
4021	Warum?\n	Why?\n
4335	-----...	-----...
8005	Unterernährung und Hunger\n	Malnutrition and Hunger\n
8008	Übertragbare Krankheiten\n	Communicable Diseases\n
8011	Übertragbare Krankheiten\n	Communicable Diseases\n
...	...	...
1920168	Wir kommen nun zur Abstimmung.\n	We shall now proceed to the vote.\n
1920206	Unterbrechung der Sitzungsperiode\n	Adjournment of the session\n
1920207	Ich erkläre die Sitzungsperiode des Europäisch...	I declare the session of the European Parliame...
592353	Da die Fettleibigkeit bei Kindern auch eng mit...	As childhood obesity is also strongly linked t...
1812042	Dieses Recht schließt die Meinungsfreiheit und...	This right shall include freedom to hold opini...

45016 rows × 2 columns

## DROPPING DUPLICATES

```
In [20]: final_data_df = final_data_df.drop_duplicates()
```

```
In [21]: print(final_data_df.count())
```

```
German      4476170
English     4476311
dtype: int64
```

# Summary of the Approach to EDA and Pre-processing:

Our agenda is to convert one language into another language. we have the text data required for the process, so we not doing any visualization feature here. As part of the preprocessing, we have followed the following steps.

- Merge, Load and examined the data
- Cleansing and processing the data
- Tokenization
- Padding

## Cleansing and processing the data:

- The data has been converted into lower case.
- Duplicating records have been dropped
- Punctuation has been replaced by empty string.
- Missing values are identified
- Special characters like '\n' has been removed

# Converting data into lower case:

```
In [145]: # Converting the data to Lower case the sentence
final_data_df['English'] = final_data_df['English'].str.lower()
final_data_df['German'] = final_data_df['German'].str.lower()
```

```
In [146]: final_data_df.head()
```

Out[146]:

	German	English
0	steigt gold auf 10000 dollar	10000 gold
1	san francisco – es war noch nie leicht ein rat...	san francisco – it has never been easy to have...
2	in letzter zeit allerdings ist dies schwierige...	lately with gold prices up more than 300 over ...
3	erst letzten dezember verfassten meine kollege...	just last december fellow economists martin fe...
4	und es kam wie es kommen musste	wouldn't you know it

```
In [147]: final_data_df.tail()
```

Out[147]:

	German	English
2399118	schon früh erkennt er dass erfolgreiche arbeit...	reiner zeising has been at home in the world o...
2399119	in seinem unternehmen werden durch diese einzi...	right from the outset he recognised that succe...
2399120	individualität ist nur eine der vielen stärken...	acting on this insight as long ago as 1988 he ...
2399121	gerhard menz feldwebel ungarn 1945 dx06exkl	kit carson 101st airborne normandy 1944 chexcl
2399122	sie ist die ganze zeit im umbau aber ich glaub...	its still under construction but i hope that a...

## Punctuation has been replaced by empty string:

In [22]: `import string`

```
final_data_df['German'] = final_data_df['German'].str.replace('{}'.format(string.punctuation), '')
final_data_df['English'] = final_data_df['English'].str.replace('{}'.format(string.punctuation), '')
final_data_df.head()
```

C:\Users\navit\AppData\Local\Temp\ipykernel\_4504\1878863163.py:3: FutureWarning: The default value of regex will change from True to False in a future version.

```
final_data_df['German'] = final_data_df['German'].str.replace('{}'.format(string.punctuation), '')
```

C:\Users\navit\AppData\Local\Temp\ipykernel\_4504\1878863163.py:4: FutureWarning: The default value of regex will change from True to False in a future version.

```
final_data_df['English'] = final_data_df['English'].str.replace('{}'.format(string.punctuation), '')
```

Out[22]:

	German	English
0	Steigt Gold auf 10000 Dollar\n	10000 Gold\n
1	SAN FRANCISCO – Es war noch nie leicht ein rat...	SAN FRANCISCO – It has never been easy to have...
2	In letzter Zeit allerdings ist dies schwierige...	Lately with gold prices up more than 300 over ...
3	Erst letzten Dezember verfassten meine Kollege...	Just last December fellow economists Martin Fe...
4	Und es kam wie es kommen musste\n	Wouldn't you know it\n

## Missing values are identified :

```
In [27]: final_data = final_data_df[final_data_df.isna().any(axis=1)]
```

```
In [28]: print(final_data)
```

	German	English
201854	NaN	Last year UN Secretary General Ban Kimoon and ...
201855	NaN	Regrettably the UN failed to follow up with in...
201856	NaN	The UN would say it must be invited into a cou...
201857	NaN	But it is also true that with adequate politic...
201858	NaN	The UN should have put Mugabe on the defensive...
...	...	...
201990	NaN	Their achievement remains one of the greatest ...
201991	NaN	At the same time Zuma's revolutionary generati...
201992	NaN	In a region that reveres the elderly Zuma's at...
201993	NaN	Three in ten South Africans are younger than 1...
201994	NaN	Somehow Zuma must find a way to honor his own ...

```
[141 rows x 2 columns]
```

Special characters like '\n' has been removed :

```
In [23]: #replace '\n' with empty string
final_data_df.replace('\n', '', regex=True, inplace=True)
```

## Tokenization :

```
In [31]: def tokenize(sentences):
# Create tokenizer
text_tokenizer = Tokenizer()
# Fit texts
text_tokenizer.fit_on_texts(sentences)
return text_tokenizer.texts_to_sequences(sentences), text_tokenizer
```

```
In [75]: # Tokenize words

eng_text_tokenized, eng_text_tokenizer = tokenize(english_sentences)
german_text_tokenized, german_text_tokenizer = tokenize(german_sentences)

print('Maximum length english sentence: {}'.format(len(max(eng_text_tokenized, key=len))))
print('Maximum length german sentence: {}'.format(len(max(german_text_tokenized, key=len))))

# Check Language Length
english_vocab = len(eng_text_tokenizer.word_index) + 1
german_vocab = len(german_text_tokenizer.word_index) + 1

print("English vocabulary is of {} unique words".format(english_vocab))
print("german vocabulary is of {} unique words".format(german_vocab))

Maximum length english sentence: 93
Maximum length german sentence: 89
English vocabulary is of 9163 unique words
german vocabulary is of 12931 unique words
```

## Padding :

When the sequence of word id's feed into the model, each sequence should have the same length. Padding is added to the shorter sequence with the max length

```
In [104]: from tensorflow.keras.preprocessing.sequence import pad_sequences

def padding(sentences,length=None):
    return pad_sequences(sentences, padding="post", truncating="post", maxlen=length)
```

## DESIGN :

### Deciding Models and Model Building :

We build the model using LSTM and RNN, by using these techniques we handled designing, training and testing.

Tokenized index has been converted into texts

```
#define the function to convert the numerals to text
def logits_to_text(logits, tokenizer):
    """
    Turn logits from a neural network into text using the tokenizer
    :param logits: Logits from a neural network
    :param tokenizer: Keras Tokenizer fit on the labels
    :return: String that represents the text of the logits
    """
    index_to_words = {id: word for word, id in tokenizer.word_index.items()}
    index_to_words[0] = '<PAD>'

    return ' '.join([index_to_words[prediction] for prediction in np.argmax(logits, 1)])

print(`logits_to_text` function loaded.)

`logits_to_text` function loaded.
```

Execution of the tokenize function. This creates the maximum length of the sequence and the number of words in the vocabulary list

```
# Tokenize words

eng_text_tokenized, eng_text_tokenizer = tokenize(english_sentences)
german_text_tokenized, german_text_tokenizer = tokenize(german_sentences)

print('Maximum length english sentence: {}'.format(len(max(eng_text_tokenized, key=len))))
print('Maximum length german sentence: {}'.format(len(max(german_text_tokenized, key=len))))

# Check Language Length
english_vocab = len(eng_text_tokenizer.word_index) + 1
german_vocab = len(german_text_tokenizer.word_index) + 1

print("English vocabulary is of {} unique words".format(english_vocab))
print("german vocabulary is of {} unique words".format(german_vocab))
```



## Output of the tokenized data

```
Maximum length english sentence: 93
Maximum length german sentence: 89
English vocabulary is of 9163 unique words
german vocabulary is of 12931 unique words
```

Data needs to be reshaped in order to execute the model. The model would need 3-dimensional data. Shape of the data has been converted to fit the model

```
# Try experimenting with the size of that dataset
max_english_len = int(len(max(eng_text_tokenized, key=len)))
max_german_len = int(len(max(german_text_tokenized, key=len)))

eng_pad_sentence = pad_sequences(eng_text_tokenized, max_english_len, padding = "post")
german_pad_sentence = pad_sequences(german_text_tokenized, max_german_len, padding = "post")

# Reshape data
eng_pad_sentence = eng_pad_sentence.reshape(*eng_pad_sentence.shape, 1)
german_pad_sentence = german_pad_sentence.reshape(*german_pad_sentence.shape, 1)
```

## Training the model : LSTM model

LOSS TREND							
	LSTM	2.2425	2.2409	2.2432	2.2454	2.2476	2.2519
Accuracies							
	LSTM	0.7095	0.752	0.7535	0.7541	0.7547	0.7548
BLEU SCORE							
	LSTM	5.57	5.039	5.32	5.79	5.57	5.79

Implementing a simple LSTM model using the encoder and decoder methods

```
#parameters for LSTM model
input_sequence = Input(shape=(max_english_len,))
embedding = Embedding(input_dim=english_vocab, output_dim=128,)(input_sequence)
encoder = LSTM(64, return_sequences=False)(embedding)
r_vec = RepeatVector(max_german_len)(encoder)
decoder = LSTM(64, return_sequences=True, dropout=0.2)(r_vec)
logits = TimeDistributed(Dense(german_vocab))(decoder)
```

Softmax activation has been implemented with accuracy metric. Summary of the model is printed

```
#Build the model
enc_dec_model = Model(input_sequence, Activation('softmax')(logits))
enc_dec_model.compile(loss=sparse_categorical_crossentropy,
                      optimizer=Adam(1e-3),
                      metrics=['accuracy'])
enc_dec_model.summary()
```

Summary model of the LSTM model is given below

Model: "model"

Layer (type)	Output Shape	Param #
=====		
input_1 (InputLayer)	[(None, 93)]	0
embedding (Embedding)	(None, 93, 128)	1172864
lstm (LSTM)	(None, 64)	49408
repeat_vector (RepeatVector)	(None, 89, 64)	0
lstm_1 (LSTM)	(None, 89, 64)	33024
time_distributed_10 (TimeDistributed)	(None, 89, 12931)	840515
activation (Activation)	(None, 89, 12931)	0
=====		

Total params: 2095811 (7.99 MB)  
Trainable params: 2095811 (7.99 MB)  
Non-trainable params: 0 (0.00 Byte)

Fit the model with early stopping based on the test and train data split with a epoch of 50 and batch size 128

```
#fit the model
from tensorflow.keras.callbacks import EarlyStopping, ReduceLROnPlateau, ModelCheckpoint
es = EarlyStopping(monitor='val_accuracy',mode='max',verbose=1,patience=40)
rl = ReduceLROnPlateau(monitor='val_accuracy',mode='max',verbose=1,patience=5,factor=0.1,min_lr=0.001)
mc = ModelCheckpoint('checkpoint/',monitor='val_accuracy',verbose=1,mode='max',save_best_only=True)

r = enc_dec_model.fit(X_train,y_train,
                      validation_data=(X_test,y_test),
                      epochs=50, batch_size=128, callbacks=[es,rl,mc])
```


## Execution output with the above metrics

```
Epoch 1/50
18/18 [=====] - ETA: 0s - loss: 9.1490 - accuracy: 0.7095
Epoch 1: val_accuracy improved from -inf to 0.74087, saving model to checkpoint\
INFO:tensorflow:Assets written to: checkpoint\assets

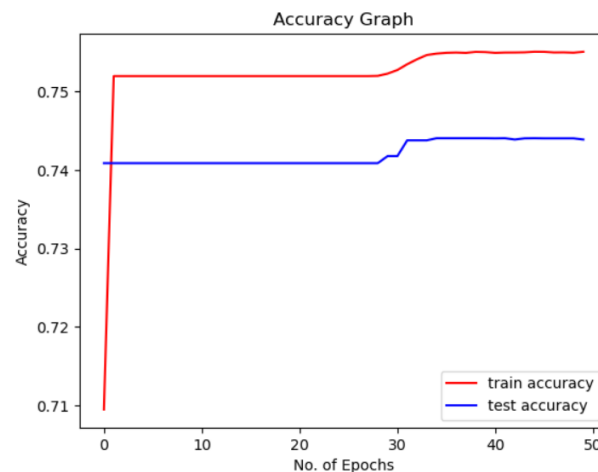
INFO:tensorflow:Assets written to: checkpoint\assets

18/18 [=====] - 65s 3s/step - loss: 9.1490 - accuracy: 0.7095 - val_loss: 8.3199 - val_accuracy: 0.7
409 - lr: 0.0010
Epoch 2/50
18/18 [=====] - ETA: 0s - loss: 7.2191 - accuracy: 0.7520
Epoch 2: val_accuracy did not improve from 0.74087
18/18 [=====] - 39s 2s/step - loss: 7.2191 - accuracy: 0.7520 - val_loss: 5.9471 - val_accuracy: 0.7
409 - lr: 0.0010
Epoch 3/50
18/18 [=====] - ETA: 0s - loss: 4.9269 - accuracy: 0.7520
Epoch 3: val_accuracy did not improve from 0.74087
18/18 [=====] - 38s 2s/step - loss: 4.9269 - accuracy: 0.7520 - val_loss: 3.9160 - val_accuracy: 0.7
409 - lr: 0.0010
Epoch 4/50
```

## Loss graph



## Accuracy graph of the models



## Testing the model : SimpleRNN model

Logit function and the model prediction of the Simple RNN model

LOSS TREND							
	RNN	2936140	2936107	2936105	2936103	2936105	2936104

```
#define and print the prediction for LSTM model
def logits_to_sentence(logits, tokenizer):

    index_to_words = {idx: word for word, idx in tokenizer.word_index.items()}
    index_to_words[0] = '<empty>'

    return ' '.join([index_to_words[prediction] for prediction in np.argmax(logits, 1)])

#index = 10
for index in range(200, 216):
    print("The english sentence is: {}".format(english_sentences[index]))
    print("The german sentence is: {}".format(german_sentences[index]))
    print('The predicted sentence is :')
    print(logits_to_sentence(enc_dec_model.predict(eng_pad_sentence[index:index+1])[0], german_text_tokenizer))
    print("BLEU Score:", sentence_bleu(german_sentences[index], logits_to_sentence(enc_dec_model.predict(eng_pad_sentence[index:index+1])[0], german_text_tokenizer)))
```

Predict the simple RNN models by calling the logits function to convert the index to text

## Testing the model : LSTM model

```
#define and print the prediction for LSTM model
def logits_to_sentence(logits, tokenizer):

    index_to_words = {idx: word for word, idx in tokenizer.word_index.items()}
    index_to_words[0] = '<empty>'

    return ' '.join([index_to_words[prediction] for prediction in np.argmax(logits, 1)])

#index = 10
for index in range(200, 216):
    print("The english sentence is: {}".format(english_sentences[index]))
    print("The german sentence is: {}".format(german_sentences[index]))
    print('The predicted sentence is :')
    print(logits_to_sentence(enc_dec_model.predict(eng_pad_sentence[index:index+1])[0], german_text_tokenizer))
    print("BLEU Score:", sentence_bleu(german_sentences[index], logits_to_sentence(enc_dec_model.predict(eng_pad_sentence[index:index+1])[0], german_text_tokenizer)))
```

## Output of the predicted model

[illegible]